

CMOR 421 Homework 2: OpenMP

Hubert King

March 25th, 2024

1 Blocked Matrix-Matrix Multiplication using OpenMP

1.1 Instructions

Optimize blocked matmul using OpenMP. You may run this on either your own laptop or on NOTS. Compare `#omp parallel for` with `#omp parallel for collapse(2)`. Report the parallel efficiency as you increase the number of threads.

1.2 Solution

Timing data and parallel efficiency for serial and parallel implementations of blocked matrix-matrix multiplication using `#omp parallel for` with `#omp parallel for collapse(2)` are summarized below. As the number of threads increased, the parallel efficiency decreased, likely due to the overhead incurred from spawning new threads. The collapse directive improved efficiency, likely due to more efficient work load balancing as the number of threads increased. Since the outer loop's size is determined by the block size, there were more threads than iterations when *threads* > 8. This is where the difference in efficiencies is greatest; the collapse directive ensures that the number of iterations in the loop is greater than the threads available.

Number of Threads	8	12	16	20	24
Serial (s)	2.51	2.51	2.51	2.52	2.51
Nested Parallel (s)	0.323	0.338	0.337	0.360	0.363
Collapse Parallel (s)	0.305	0.233	0.232	0.262	0.319

Table 1: Blocked matrix multiplication execution time for serial and parallel implementations using nested parallelism and collapse directive. Timings reflect the total execution time for five trials. For all trials, block size was $n = 8$




Figure 1: Parallel efficiency plot for parallel back solve algorithm with static scheduling. Five trials were run and eight threads were used.

2 Parallel Back Solve with Unit Upper Triangular Matrix

2.1 Intructions

Parallel back solve with a unit upper triangular matrix. Write out the back solve algorithm explicitly and identify parallelism. Provide pseudocode, and implement your algorithm in C++ using OpenMP. For a sufficiently large matrix size, test your implementation: Does static or dynamic scheduling work better? Why? Test your implementation for several matrix sizes and report the parallel efficiency as a function of matrix size. Do you expect parallel efficiency to be better or worse if the lower triangular matrix is only bi-diagonal?

2.2 Solution

Given a system $Ax = b$, where A is an $n \times n$ unit upper triangular matrix with any upper entries, and b is the right-hand side vector, the matrix A and vector b can be represented as:

$$A = \begin{bmatrix} 1 & a_{12} & a_{13} & \cdots & a_{1n} \\ 0 & 1 & a_{23} & \cdots & a_{2n} \\ 0 & 0 & 1 & \ddots & \vdots \\ \vdots & \vdots & \ddots & \ddots & a_{(n-1)n} \\ 0 & 0 & \cdots & 0 & 1 \end{bmatrix}, \quad b = \begin{bmatrix} b_1 \\ b_2 \\ \vdots \\ b_n \end{bmatrix}$$

The equations derived from this system are for $i = 1, \dots, n$:

$$x_i + \sum_{j=i+1}^n a_{ij}x_j = b_i$$

This system can be solved using the back solve algorithm below. This approach iterates through rows in the outer loop and columns in the inner loop.

Algorithm 1 Back Solve with Unit Upper Triangular Matrix

Input: A (an $n \times n$ unit upper triangular matrix), b (vector in \mathbb{R}^n), n (dimension of A)
Output: x (vector in \mathbb{R}^n)
for $i = 1, \dots, n$ **do**
 $x_i = b_i$
end for
for $j = n, \dots, 1$ **do**
 for $i = 1, \dots, j$ **do**
 Compute $x_i = x_i - A_{ij} \cdot x_j$
 end for
end for
return x

Algorithm 2 uses a parallel for-loop to assign values of b to x , and another parallel for-loop to adjust x as the inner loop of the nested for-loop iterates through rows of A .

Algorithm 2 Back Solve with Unit Upper Triangular Matrix

Input: A (an $n \times n$ unit upper triangular matrix), b (vector in \mathbb{R}^n), n (dimension of A)
Output: x (vector in \mathbb{R}^n)
Parallel for-loop: distribute for-loop across threads
for $i = 1, \dots, n$ **do**
 $x_i = b_i$
end for
for $j = n, \dots, 1$ **do**
 Parallel for-loop: distribute for-loop across threads
 for $i = 1, \dots, j$ **do**
 Compute $x_i = x_i - A_{ij} \cdot x_j$
 end for
end for
return x

Timing Data and parallel efficiency are summarized below. Each algorithm was run for five trials, with eight threads. The effects of parallelization were observed with $n > 3000$. Static scheduling had better effects on execution time than dynamic scheduling. This could be explained by the extra overhead incurred by dynamic scheduling. While operations per iteration in the inner for-loop do vary, meaning dynamic scheduling should improve performance, the benefit may not be enough to offset the extra overhead experienced.

For a bi-diagonal unit upper triangular matrix, parallel efficiency would likely be greater, since the inner loop performs exactly two operations per loop. Thus, the spread of work across threads is maximally efficient with the least possible downtime. Thus, static scheduling could be used with no risk of load imbalance. If we consider the grocery checkout analogy presented in class, this means that all cashiers have two customers at a time, and would get new customers simultaneously with no downtime.

Matrix Size	3000	6000	9000	12000	15000	18000	21000	24000	27000	30000
Serial (s)	0.130	0.662	1.49	2.82	4.29	6.30	9.73	17.28	14.5	18.3
Static (s)	0.114	0.236	0.414	0.87	0.98	1.43	2.29	3.42	2.65	3.37
Dynamic (s)	1.93	8.39	18.3	-	-	-	-	-	-	-

Table 2: Back Solve algorithm execution time for serial implementation and parallel implementation with static and dynamic scheduling. All implementations used eight threads. Dynamic timings were omitted past $n = 9000$, and values reflect the total execution time for five trials.

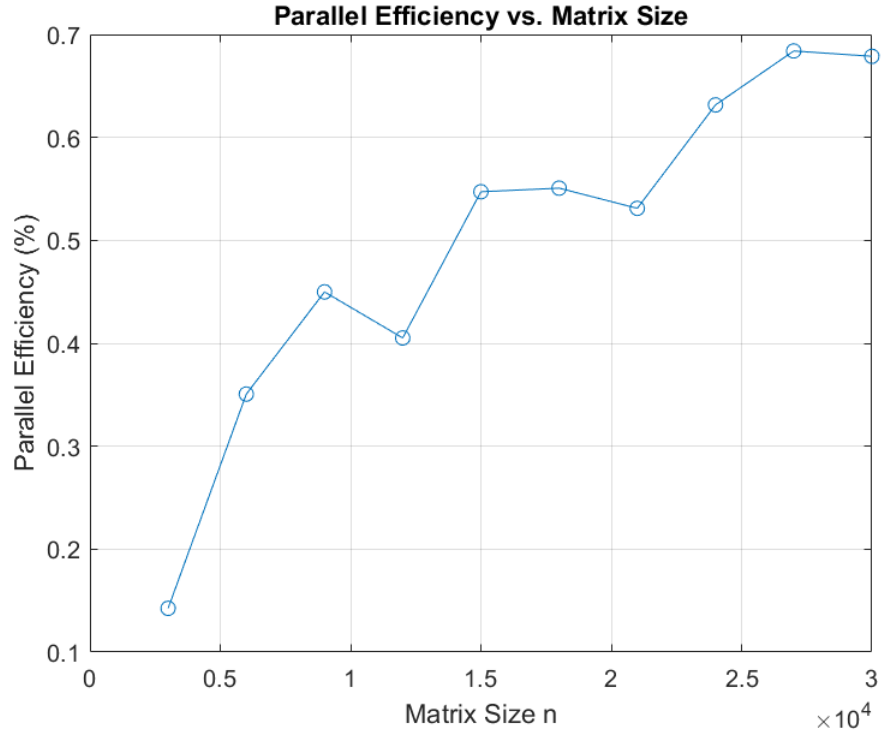


Figure 2: Parallel efficiency plot for parallel back solve algorithm with static scheduling. Five trials were run and eight threads were used.

3 Build and Run Instructions

3.1 Prerequisites

Ensure the following requirements are met:

- The instructions are intended for users on the NOTS system.
- GCC version 13.1.0 is necessary for successful compilation and execution of the rmatmul.cpp.

3.2 Build and Run

Access NOTS via a login node and load the necessary modules:

```
module load GCCcore/13.1.0
```

Verify that the module is loaded correctly and the correct version of GCC is being used:

```
gcc --version
```

Next, compile the program with the following command:

```
g++ -fopenmp -o main -Iinclude main.cpp src/functions.cpp
```

After successful compilation, the program can be tested by running the following command in the interactive node, where dimension is replaced by the desired matrix size. This size will be used for both blocked matrix-matrix multiplication and back-solve algorithms.

```
./main <dimension>
```

To see the effects of multi-threading, access an interactive node and specify that 24 CPUs are to be allocated:

```
srun --partition=interactive --pty --export=ALL --cpus-per-task=24 --time=00:30:00 /bin/bash
```

Now, run the executable as before:

```
./main <dimension>
```