

CMOR 421 Homework 3: MPI

Hubert King

April 9th, 2024

1 Simplified SUMMA Algorithm

For our simplified SUMMA, we use the following approach. Algorithm 1 provides a more detailed outline.

- Scatter partitions of A and B to the processes, stored as A^{local} and B^{local} .
- Set $C^{local} \leftarrow 0$.
- Set $blocksize \leftarrow n \div \sqrt{s}$.
- Allocate buffers A^{buffer} and B^{buffer} .
- Initialize row and column communicators.
- In the main computational loop, all processes perform p iterations of the following:
 - Broadcast A^{local} across its row.
 - Broadcast B^{local} across its column.
 - Receive row broadcast into A^{buffer} .
 - Receive column broadcast into B^{buffer} .
 - Accumulate $C^{local} \leftarrow A^{buffer} \times B^{buffer}$.
- Gather blocks of C from all processes onto root process.

2 Cannon's Algorithm

For Cannon's algorithm, we use the following approach, keeping the same *blocksize* from SUMMA and procedure for the initial scatter of the partitions of A , B , and C .

- For each row i of A , perform a left-rotating shift of the partitions by i position.
- For each row j of B , perform a upper-rotating shift of the partitions by j position.
- In the main computational loop, all processes perform p iterations of the following:
 - Accumulate $C^{local} \leftarrow A^{local} \times B^{local}$
 - Perform left-circular shift of A on all rows by 1 position.
 - Perform upward-circular shift of B on all columns by 1 position.
- Gather blocks of C from all processes onto the root process.

Algorithm 1 SUMMA Algorithm Pseudocode

```
1: Input:  $A, B, rank, size$ 
2: Output:  $C$ 
3:  $blocksize \leftarrow n/\sqrt{p}$ 
4: if  $rank = 0$  then
5:   for  $k = p - 1$  to  $0$  do
6:     for  $i = 1$  to  $blocksize$  do
7:       for  $j = 1$  to  $blocksize$  do
8:          $A_{ij}^{local} \leftarrow A_{ij}$ 
9:          $B_{ij}^{local} \leftarrow B_{ij}$ 
10:         $C_{ij}^{local} \leftarrow 0$ 
11:      end for
12:    end for
13:    if  $k > 0$  then
14:      Send  $A^{local}, B^{local}, C^{local}$  to process  $k$ .
15:    end if
16:  end for
17: end if
18: if  $rank > 0$  then
19:   Receive  $A^{local}, B^{local}, C^{local}$  from  $rank = 0$ .
20: end if
21: for  $k = 0$  to  $\sqrt{p} - 1$  do
22:   for  $i = 1$  to  $\sqrt{p}$  do
23:     Process holding block  $A(i, k)$  broadcasts to its row.
24:   end for
25:   for  $j = 1$  to  $\sqrt{p}$  do
26:     Process holding block  $B(k, j)$  broadcasts to its column.
27:   end for
28:   Receive block  $A(i, k)$  into  $A^{buffer}$ .
29:   Receive block  $B(k, j)$  into  $B^{buffer}$ .
30:   for  $i = 1$  to  $blocksize$  do
31:     for  $j = 1$  to  $blocksize$  do
32:        $C_{ij}^{local} \leftarrow C_{ij}^{local} + A_{ij}^{buffer} \cdot B_{ij}^{buffer}$ 
33:     end for
34:   end for
35: end for
36: if  $rank > 0$  then
37:   Send  $C^{local}$  to  $rank = 0$ .
38: end if
39: if  $rank = 0$  then
40:    $C \leftarrow C^{local}$ 
41:   for  $k = size$  to  $0$  do
42:     Receive  $C^{local}$  from  $rank = k$ 
43:      $C \leftarrow C^{local}$ 
44:   end for
45: end if
46: return  $C$ 
```

3 Miscellaneous Details

To generate a random matrices for testing, we utilize the random number generation engine provided by the C++ standard library. We perform a correctness check by comparing product C from the serial and parallel algorithms element-wise, with a tolerance of $1e-9$. Timing results are measured in milliseconds.

4 Build and Run Instructions

Access NOTS via a login node and load the necessary modules:

```
module load GCCcore/13.2.0
module load OpenMPI
```

Verify that the module is loaded correctly and the correct version of GCC is being used:

```
gcc --version
mpic++ --version
```

Next, compile the drivers with the following command:

```
mpic++ -o summa -Iinclude summa.cpp src/functions.cpp
mpic++ -o cannon -Iinclude cannon.cpp src/functions.cpp
```

After successful compilation, the programs can be tested by running the following command in the login node, where dimension is replaced by the desired matrix size.

```
mpirun -n <processors> cannon <dimension>
mpirun -n <processors> summa <dimension>
```

To run, we use the following script, named job.slurm, which requests resources and runs the program on NOTS:

```
#!/bin/bash
#SBATCH --job-name=CMOR-421-521
#SBATCH --partition=scavenge
#SBATCH --reservation=CMOR421
#SBATCH --ntasks=<requested-processors>
#SBATCH --mem-per-cpu=1G
#SBATCH --time=00:30:00
echo "My job ran on:"
echo $SLURM_NODELIST
srun -n <processors> summa <dimension>
srun -n <processors> cannon <dimension>
```

Submit the job with the following command:

```
sbatch job.slurm
```

After job completion, view the output with the following command:

```
cat slurm-<job-number>.out
```

Sample output:

```
My job ran on:
bc8u27n1
Matrix size n = 1024
Serial elapsed time = 8132.66
Elapsed time = 2.37754
```

Serial product and Cannon's product are equal to machine precision.
Matrix size $n = 1024$
Serial elapsed time = 8123.66
Elapsed time = 2.37487
Serial product and SUMMA product are equal to machine precision.