

# CMOR 421/521:

## Order Dependency and Synchronization

Date: 2/13/2024

This Week:

**Tu: Order Dependency and Synchronization**

Th: Load Balancing, Reductions, and Thread Mapping

**Note:** I will be available after class today and Thursday to help with code issues.

Tu	Th	
		1
		2
		3
		4
		5
		6
		7
		8
		9
		10
		11
		12
		13
		14
		15

# Today

- Recap of some things from last Tuesday
- Race Conditions
- Order Dependency
- Synchronization

# How to OpenMP

- OpenMP uses precompiler *directives* to define parallel regions
- Directives can be modified using clauses

```
#pragma omp <directive> <clauses> <myb other stuff>
{
    // Code you want to run in parallel
}
```

# Using OpenMP: Parallel Blocks


- The simplest OpenMP directive: **parallel**

```
// Parallel hello world

#pragma omp parallel
{ // Code standards: allowed for OpenMP
    tid = omp_get_thread_num();
    printf("Hello from thread %d\n", tid);
}
// Note: threads are zero-indexed
```


# OpenMP: Thread Branching

- Once you enter a parallel block, the code branches to each thread



# OpenMP: Fork-Join Parallelism

- Master thread spawns a team of threads as needed.
- Parallelism added incrementally until performance goals are met, i.e., the sequential program evolves into a parallel program.



# OpenMP: Parallel Blocks Cont.


- Each thread does **everything** in the parallel block

Example: you can put a regular for-loop inside a parallel block

```
#pragma omp parallel
{
    for(int i = 0; i < n; i++) {
        A[i] += i;
    } // Every thread will do the whole loop
}
```

- **Threads execute in any order!**
  - They are also in parallel, and can interrupt each other
  - Example: printing is a mess!

# An example problem: numerical integration



Mathematically, we know that:

$$\int_0^1 \frac{4.0}{(1+x^2)} dx = \pi$$

We can approximate the integral as a sum of rectangles:

$$\sum_{i=0}^N F(x_i) \Delta x \approx \pi$$

where each rectangle has width  $\Delta x$  and height  $F(x_i) = 4/(1+x^2)$  at the middle of interval  $i$ .




# An example serial code

- See `integration.cpp`
- This is a “second order” accurate method; roughly speaking, if we want 14 digits of accuracy, we’ll need  $1e7$  function evaluations.
- This code resulted in a *race condition*

# First try at a (correct) solution:

- Ensure no two threads write to the same blocks of memory.
- Simple parallelization pattern: for 4 threads




# What about scalability?

- See `integration_omp_no_race.cpp`

# What is the issue?

- The memory hierarchy!
- Each processor has *its own L1 cache (or cache hierarchy)*
  - Cache hierarchy behaves similarly to the single processor case.
- Challenge: how do you ensure **cache coherency**?




Idealized model of shared memory for processors P1, ..., Pn

# Why poor scaling?

- “A transaction is *relevant* if it involves a cache block **currently contained in this cache**”
- Suppose you have two threads which access memory locations close to each other.
  - Thread 1 pulls a cache line into Cache 1
  - Thread 2 pulls a cache line into Cache 2

What happens if Cache 1 memory gets updated?



# How to fix false sharing?

- **Padding:** ensure memory accessed by different threads do not share the same cache line
  - Add unused entries in between each
- Note: padding may or may not result in the same speedup on all architectures.
  - Cache coherence protocol might be less expensive
  - Smaller number of threads? Compiler takes care of it?
- See `integration_omp_padding.cpp`

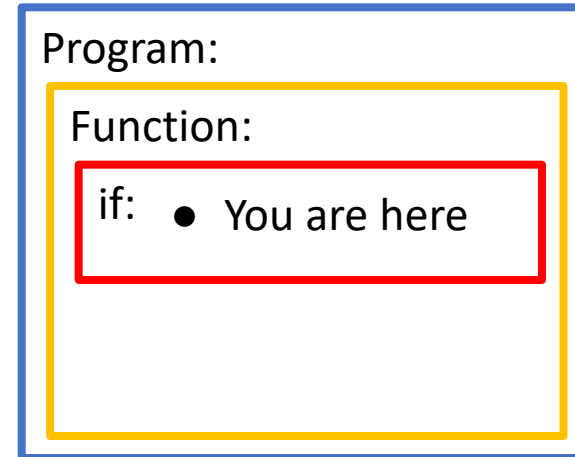
# How to fix, but less ugly?

- Padding is fairly invasive
- OpenMP provides several nice alternatives
  - Option 1: **private** variables and **critical** regions
  - Option 2: **omp for**: tries to automate worksharing


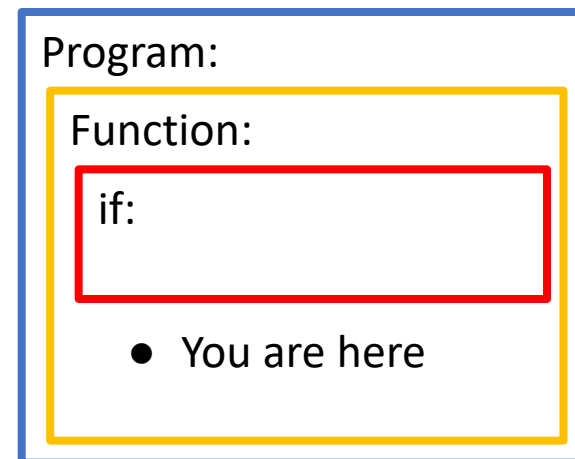

# The Stack Revisited

- The stack data structure is like a stack of cups: Last-In-First-Out (LIFO)
- Stack memory allocations operate within a single program as a stack
- Allocation and deallocation uses LIFO ordering to ensure scope

Program Scope:



Stack memory:




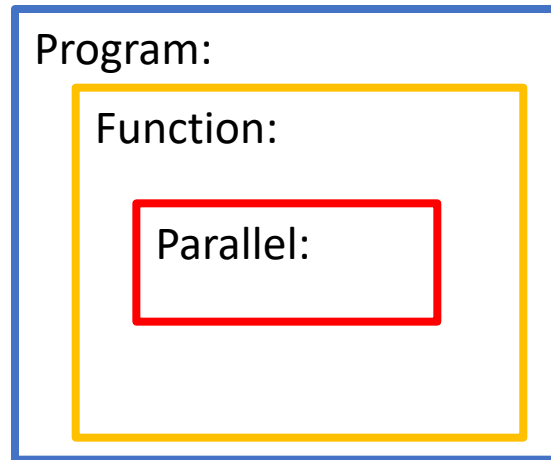


# OpenMP: Variable Scope

- The parallel block has braces around it: it scopes variables
- OpenMP uses a shared memory paradigm: which variables are shared, which are not?

# The Stack in Parallel

- Each thread has its own variables on the stack
  - Imagine being able to have multiple cups in a given layer
- LIFO structure remains, so does scope



# OpenMP: Variable Scope & Shared Memory

- **Q:** Which variables are shared, which are not?
- **A:** By default:
  - (Stack) Variables declared before the parallel block are **shared**, i.e. all threads have access to those variables
  - (Stack) Variables declared inside the block are **private**, i.e. each thread has their own version of that variable
  - We can change this using *clauses* with our directives

# Recall: How to OpenMP

- OpenMP uses precompiler directives to define parallel regions
- Directives can be modified using clauses

```
#pragma omp <directive> <clauses> <myb other stuff>
{
    // Code you want to run in parallel
}
```

# OpenMP: Scoping Clauses

- **shared:** all threads have access to the same variable
- **private:** every thread get their own instance of a variable, but it is not initialized
- **firstprivate:** same as private, but initialized
- **default:** sets the scoping for all variables that are used in the block

# OpenMP: Scoping Clauses

- **shared:** all threads have access to the same variable

```
#pragma omp parallel shared(x)
{
    int x;
    tid = omp_get_thread_num();
    do_stuff(tid, x);
}
```

- Variables which are defined outside the parallel section but used inside it are assumed to be **shared**

# OpenMP: Scoping Clauses

- **private:** every thread get their own instance, not initialized

```
int x = 4;
#pragma omp parallel private(x)
{    // x may not be 4 in here
    tid = omp_get_thread_num();
    do_stuff(tid, x);
}
```

- Variables defined inside `#pragma omp parallel` are private

# Example: integration

- **private:** every thread get their own instance, not initialized
  - We can make the “sum” accumulator a private variable
- Needs a synchronization directive: `#pragma omp critical`
- See `integration_omp_private.cpp`



# Scaling Revisited

- Weak and strong scaling laws were ideal cases
  - They only cared about parallel (p), sequential (s) portion
  - The act of parallelizing a code may change p and s!
- Think of the final loop that combines the thread-local sums

**Amdahl's (Strong)**

$$S(n) = \frac{1}{s + \frac{p}{n}}$$

# Order Dependency

- Order dependencies can be operator, algorithm, or instruction based
  - We can influence algorithm and instruction dependencies
  - **Algorithm:** Rectify *then* normalize a vector
  - **Instruction:** Modify a variable value
  - Operator dependencies are handled by the compiler
- **Order dependencies ensure correctness**
  - Order *sensitivities* can impact efficiency, e.g. locality

# Race Conditions

- A *race condition/data race* occurs when two threads try to read/write to the same variable
  - If only reads occur, there will be no issue
  - If even one thread tries to write, there can be an issue
    - Other threads may read and use a *stale* value
- Race conditions are an instruction-based violation of order dependency unique to parallelized codes

# Parallel Codes and Order Dependency

- Parallel codes can also have other problems:
  - Threads can execute in any order
  - They can start/**finish** in any order as a result
- What if you need to make sure the last (parallel) task was completed before moving to the next?
- This gives rise to the idea of *thread safety*
  - Can a given routine **guarantee** its correctness in a parallel setting?

# Thread Safety: Synchronization


- To preserve order dependency, we must introduce *synchronization*
- Synchronization is an operation between threads/their memory that restricts threads ability to move forward or enter a given block of code
  - Can be implicit or explicit
  - May involve all threads or just some
  - Example: **critical** directive from last week

# Implicit Synchronization

- The previously described commands are examples of explicit synchronization
  - i.e., synchronization explicitly requested by the user
- Parallel regions are also subject to implicit synchronization at their closing brace
- Note the number of threads may change from one parallel region to the next though

# Implicit Synchronization: Join

- Threads rejoin at the end of a parallel region
- When we enter the enclosing region, we know all threads have completed the inner region's code



# Implicit Synchronization

- Parallel regions are also subject to implicit synchronization at their closing brace

```
#pragma omp parallel for
for(i = 0; i < n; i++) {
    // Routine 1
} // Implicitly synchronizes here
```

```
#pragma omp parallel for
for(i = 0; i < n; i++) {
    // Routine 2
}
```

```
#pragma omp parallel
{
    #pragma omp for
    for(i = 0; i < n; i++) {
        // Routine 1
    }
    // Explicit syn. required here
    #pragma omp for
    for(i = 0; i < n; i++) {
        // Routine 2
    }
}
```



# OpenMP Synchronization Directives

OpenMP provides several directives for synchronization which do different things

## Important:

- barrier
- critical
- single

## Less important:

- atomic
- master
- ordered
- flush

# Synchronization: barrier


- **barrier** makes threads wait until all threads have reached that point in the code
- Binds to the innermost parallel region

```
#pragma omp parallel
{
    // Routine 1
    #pragma omp barrier

    // Routine 2
}
```

# Barrier and Nested Regions

- **barrier** binds the innermost parallel region
- Threads on the enclosing region are not affected



# Synchronization: **critical**

- **critical** only allows one thread to execute that block of code at a time

```
#pragma omp parallel shared(x,y)
{
    #pragma omp critical
    {
        x++;
        y++;
    }
}
```

# Synchronization: `single`

- `single` code that should only be run by one thread, but which thread does it doesn't matter

```
#pragma omp parallel
{
    ...
    // for-loop for some operation
    #pragma omp single
    {
        print("Completed work\n");
    }
}
```

# Synchronization: `atomic`

- `atomic` is similar to `critical` but can only deal with one operation

```
#pragma omp parallel shared(x)
{
    #pragma omp atomic
    x++;
}
```

# Synchronization: master

- **master** same as **single**, but now the master thread (thread 0) must be the one run the code

```
#pragma omp parallel
{
    ...
    // for-loop for some operation
    #pragma omp master
    {
        print("Completed work.\n");
    }
}
```

# Synchronization: ordered

- **ordered** specifies that a block of code in a parallel for loop must be run in order

```
#pragma omp parallel for
for(i = 0; i < n; i++) {
    // do actual work
    #pragma omp ordered
    {
        print("%d\n", i);
    }
}
```



# Synchronization: flush

- **flush** write all shared variable to memory; i.e. make memory (RAM/disk) consistent with cache values
- Can specify which vars or do all

```
#pragma omp parallel shared(x,y)
{
    // Does stuff in parallel w x,y
    #pragma omp flush
    #pragma omp flush(x)
}
```

# Effect of Synchronization

Some guiding principles:

- **A well-used system is never idle**
  - Having threads sit idle is wasteful
- **A well-used system does more work than overhead**
  - Synchronizing threads costs overhead
- **Synchronization is necessary, but expensive**
  - Don't use it unnecessarily, but don't be afraid to use it
  - Its expense also impacts how you should write code

# Synchronization and Performance

- Synchronization causes idle threads and creates overhead
  - If threads start and finish around the same time, there is less idle time
- Example problem: 23 elements, 4 threads
  - Better to have one thread do more or less than the others?

Array indices	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20	21	22
v1:	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20	21	22
v2:	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20	21	22

# Synchronization and Performance

- How much time will the different versions take with perfect parallelization?
- i.e., assume all threads:
  - Start at the same time
  - Take the same amount of time to update one element
- How much “work potential” is wasted?
  - v1: (1 idle thread)\*(1 step)
  - v2: (3 idle threads)\*(3 steps)
- This is the major motivation behind *load balancing*

v1:

t	0	1	2	3
1	0	6	12	18
2	1	7	13	19
3	2	8	14	20
4	3	9	15	21
5	4	10	16	22
6	5	11	17	

v2:

t	0	1	2	3
1	0	5	10	15
2	1	6	11	16
3	2	7	12	17
4	3	8	13	18
5	4	9	14	19
6				20
7				21
8				22

# Load Balancing with OpenMP

- How does OpenMP's `parallel-for` distributes iterations of the `for-loop`?
- Scientific computing often feature nD arrays...
  - What if there are nested `for-loops`?
  - What if there are fewer iterations than threads?
- **What if some iterations cost more than others?**

# More food for thought...

- Can synchronization cause problems?
- Consider the following code:

```
#pragma omp parallel
{
    thread_id = omp_get_thread_num();
    if thread_id == 0
        do_stuff();
    else
        #pragma omp barrier
}
```

